TalkingData原子立方体借力Druid加速海量数据的统计分析

TalkingData

移动 · 数据 · 价值

- 支持20种主流移动平台
- 40万款智能移动系统
- 每天2.5亿活跃智能设备
- 每天处理34亿会话，400亿事件
- 每天14TB数据流分发
- 支持实时统计分析与查询，部分指标需要精准统计

- 大数据量查询速度慢
- 大集群构建成本高昂
- 多维交叉计算能力低效
- 流式计算无法回朔

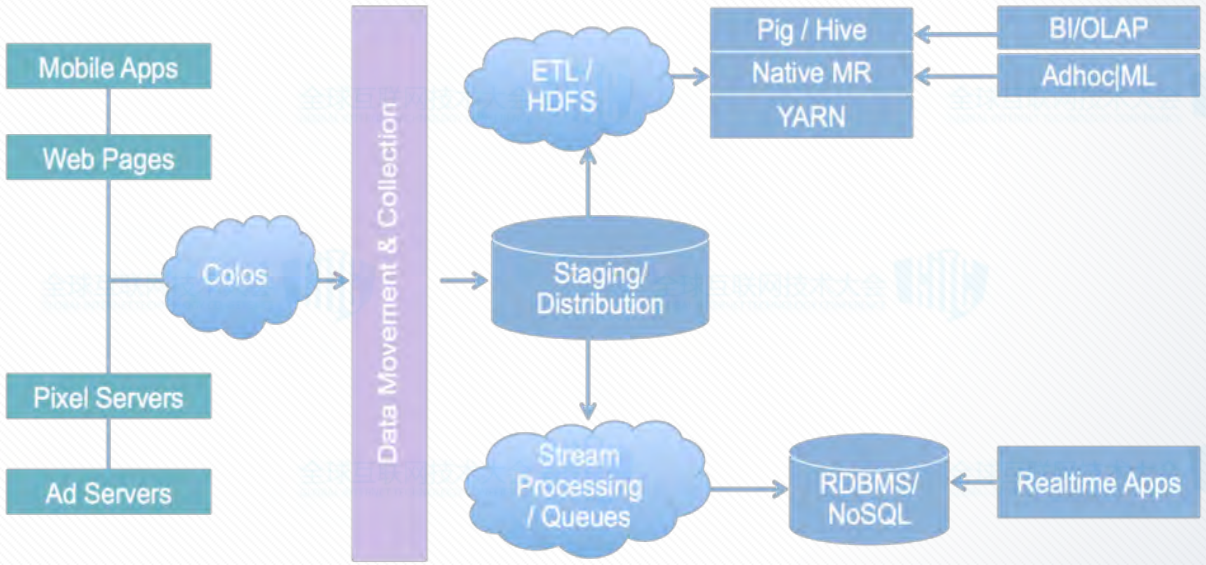

OLAP-Cube: multi-dimensional array of data

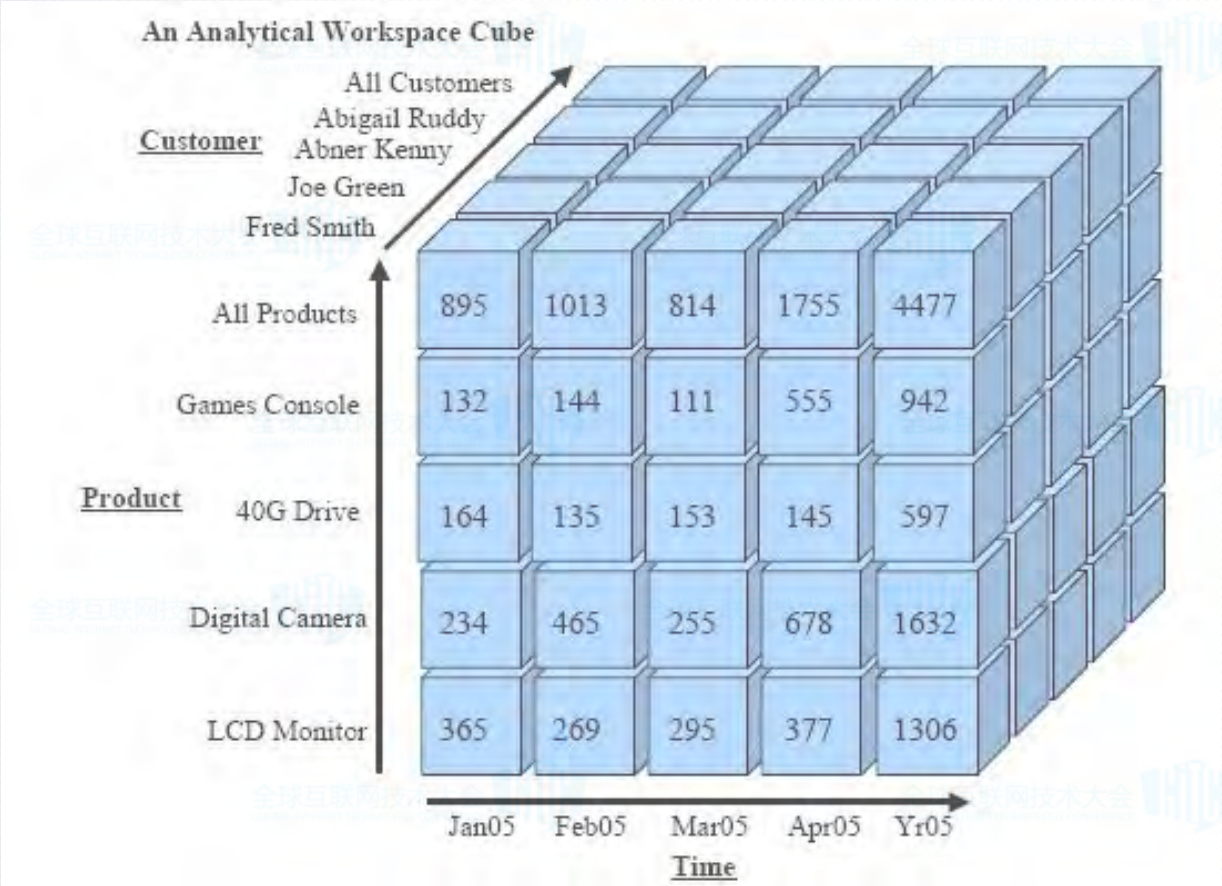

Atom OLAP Cube:

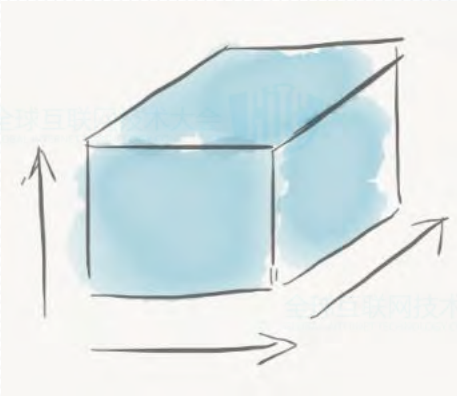

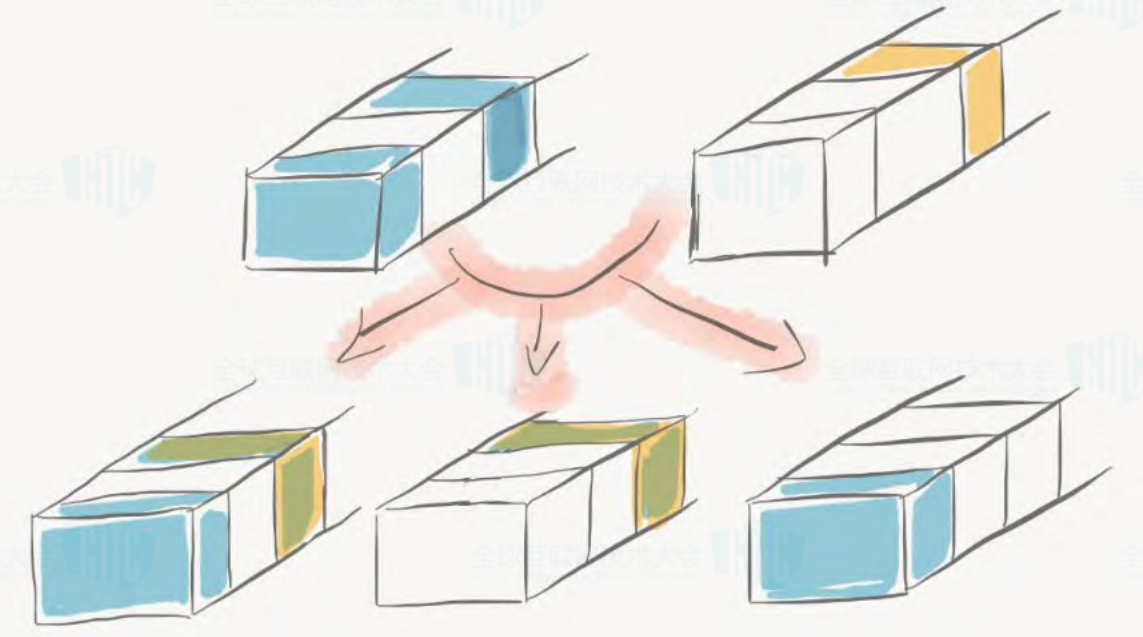

Calculation：

	∪	∩	⊕
交换	$A \cup B = B \cup A$	$A \cap B = B \cap A$	$A \oplus B = B \oplus A$
结合	$(A \cup B) \cup C = A \cup (B \cup C)$	$(A \cap B) \cap C = A \cap (B \cap C)$	$(A \oplus B) \oplus C = A \oplus (B \oplus C)$
幂等	$A \cup A = A$	$A \cap A = A$	

Fact table：

Time	Deviceid	Province	Mobile	App	Event
2015-12-22 1:00	036ca36f9c971906a97e2321ae0aeff8a	北京	Nexus S	全面枪战	充值
2015-12-22 1:30	06bd68dc66029f975c86d30e3e296d658	北京	Nexus S	全面枪战	充值
2015-12-22 2:00	02e2b5bac7ec1f9f993d48484b9fbf333	天津	iPhone 5s	滴滴打车	支付

id	Deviceid
0	036ca36f9c971906a97e2321ae0aeff8a
1	06bd68dc66029f975c86d30e3e296d658
2	02e2b5bac7ec1f9f993d48484b9fbf333

Time	id	Province	Mobile	App	Event
2015-12-22 1:00	0	北京	Nexus S	全面枪战	充值
2015-12-22 1:30	1	北京	Nexus S	全面枪战	充值
2015-12-22 2:00	2	天津	iPhone 5s	滴滴打车	支付

TD atom cube 原理

Fact table：

id	Deviceid
0	036ca36f9c971906a97e2321ae0aeff8a
1	06bd68dc66029f975c86d30e3e296d658
2	02e2b5bac7ec1f9f993d48484b9fbf333

Time	id	Province	Mobile	App	Event
2015-12-22 1:00	0	北京	Nexus S	全面枪战	充值
2015-12-22 1:30	1	北京	iPhone 5s	全面枪战	充值
2015-12-22 2:00	2	天津	Nexus S	滴滴打车	支付

TD atom cube：

Time	App	Bitmap
2015-12-22	全面枪战	0、1
2015-12-22	滴滴打车	2

Time	Mobile	Bitmap
2015-12-22	Nexus S	0、2
2015-12-22	iPhone 5s	1

Time	Dimension	Metric
Time	Provice	Bitmap
2015-12-22	北京	0、1
2015-12-22	天津	2

Time	Event	Bitmap
2015-12-22	充值	0、1
2015-12-22	支付	2

TD atom cube：

Time	App	Bitmap
2015-12-22	全面枪战	0、1（bitmap1）
2015-12-22	滴滴打车	2（bitmap2）

Time	Province	Bitmap
2015-12-22	北京	0、1（bitmap5）
2015-12-22	天津	2（bitmap6）

Time	Mobile	Bitmap
2015-12-22	Nexus S	0、2（bitmap3）
2015-12-22	iPhone 5s	1（bitmap4）

Time	Event	Bitmap
2015-12-22	充值	0、1（bitmap7）
2015-12-22	支付	2（bitmap8）

基数计算：

SELECT Distinct(Device) Where App =全面枪战 and province=北京 and time= 2015-12-22

运算转变成：Bitmap1 and bitmap5

优势：

1. 存储减少
2. 计算快
3. 支持join(这个能够很好的解决留存类型的分析)

TD Atom Cube Implementation

TD atom cube :

Bitmap+ Concise=ConciseSet

ConciseSet ： https://github.com/metamx/extendedset

RoaringBitmap

1. Storage

   Central Storage, IO exhaust.

2. Computation

   OOM

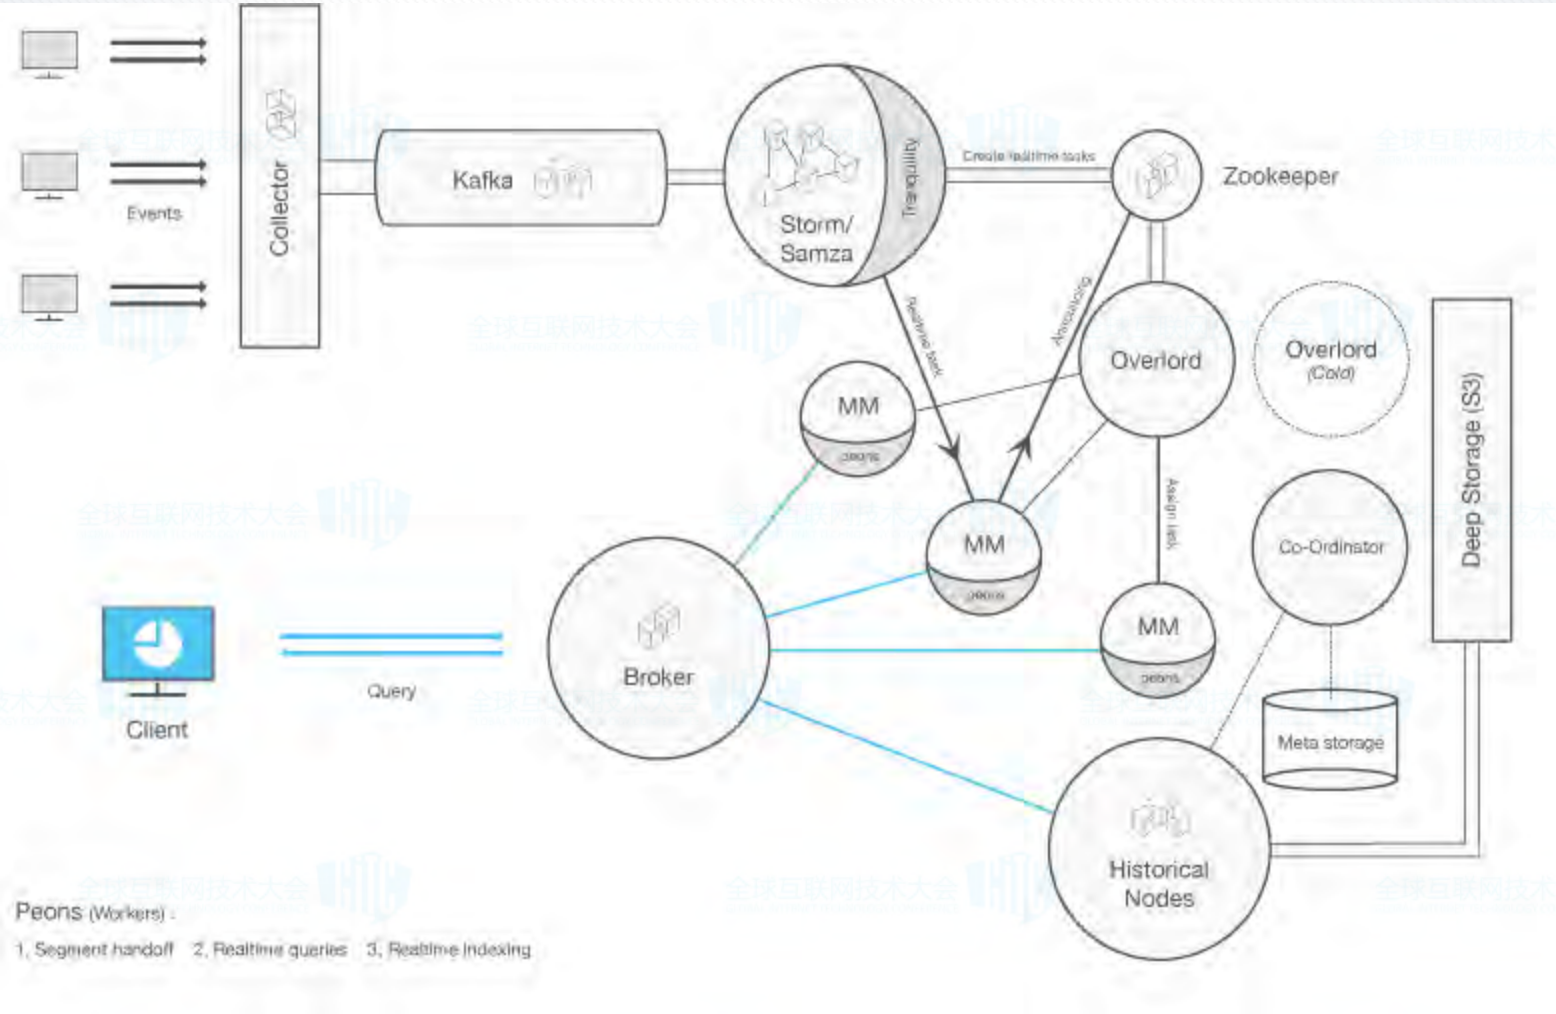

Stream/Batch ingestion

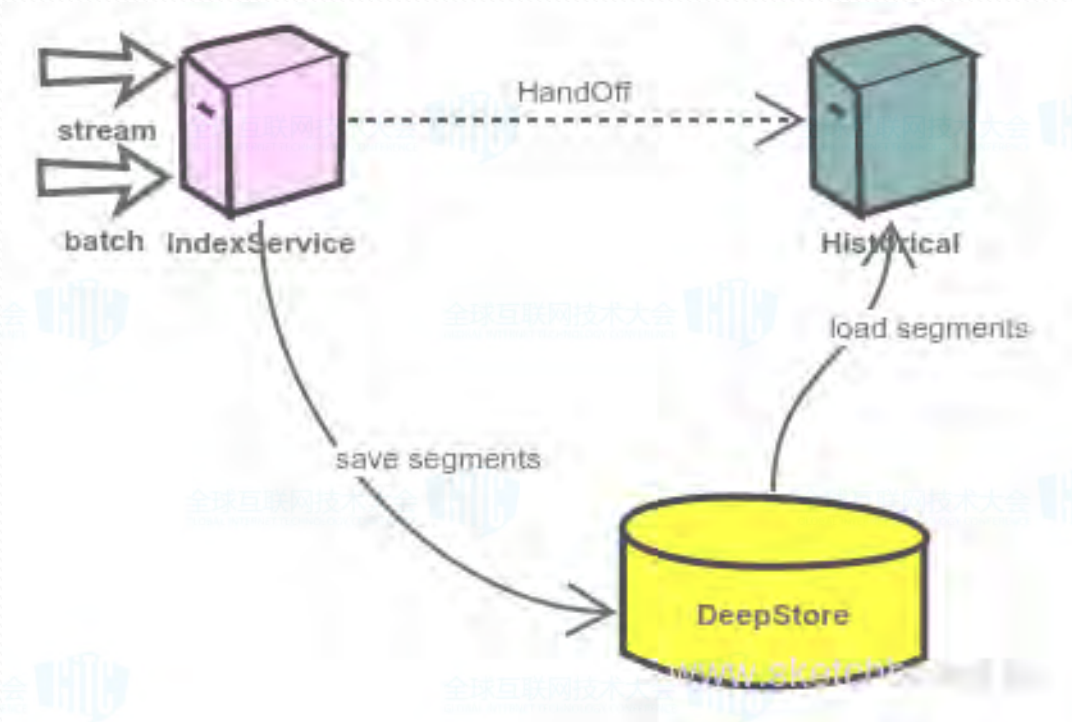

Indexing

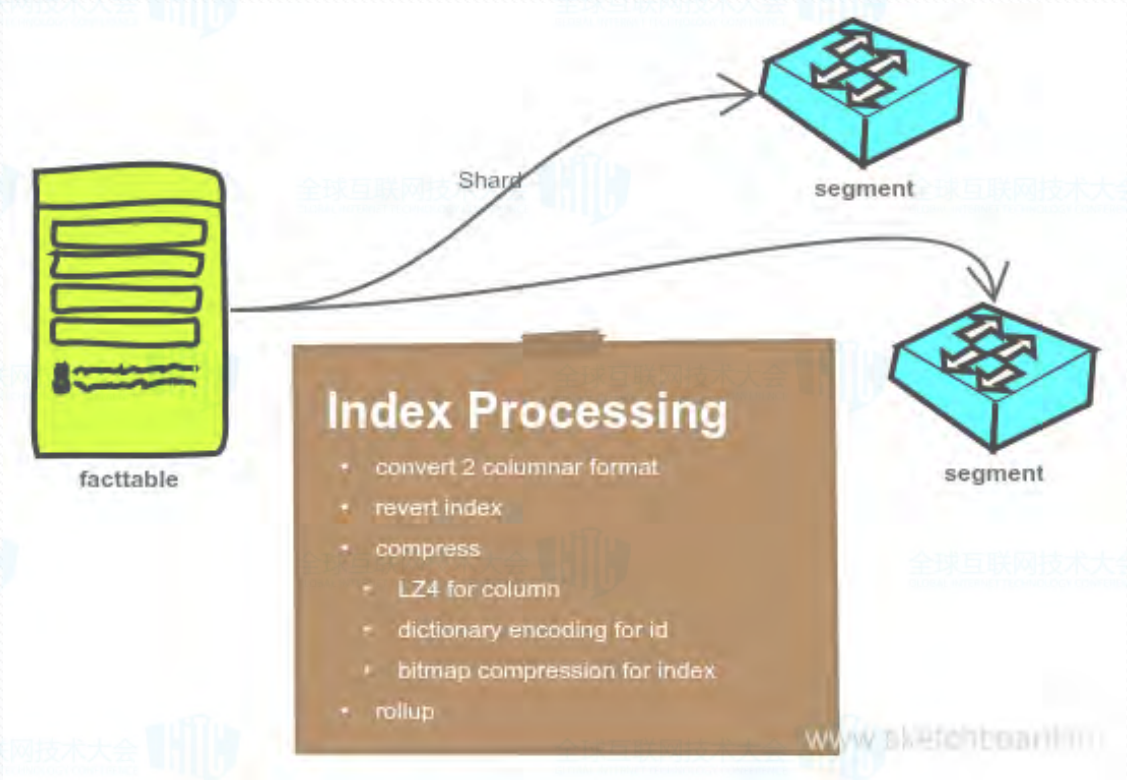

Query

- Group by
- Top N
- Timeseries
- Search
- Time boundary
- Metadata query

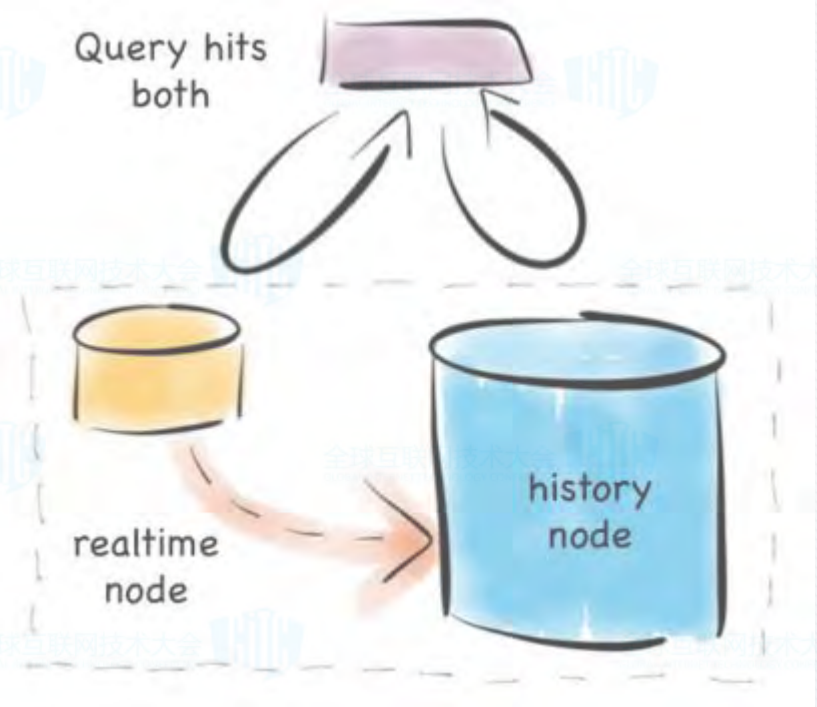

High Availability and Load Balance

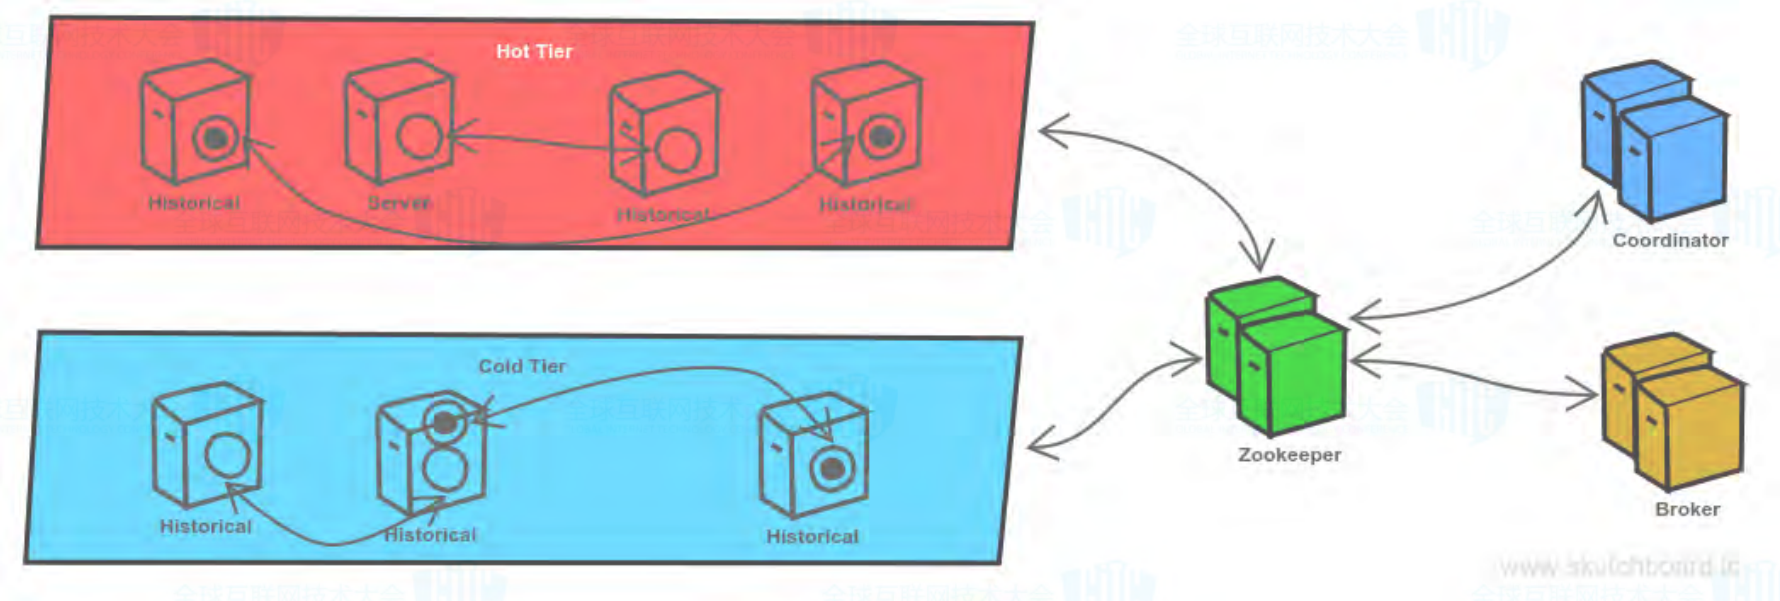

Druid's Limitation we concerned:

- Exactly cardinality calculate
  - Estimation approach by HyperUnique/Cardinality/Sketch estimat
- Join
  - limited support through “lookups”, replace one dimension value with another value

Idea：

Makes an unique id as metric, saved in Atomcube(bitmap) for each row.

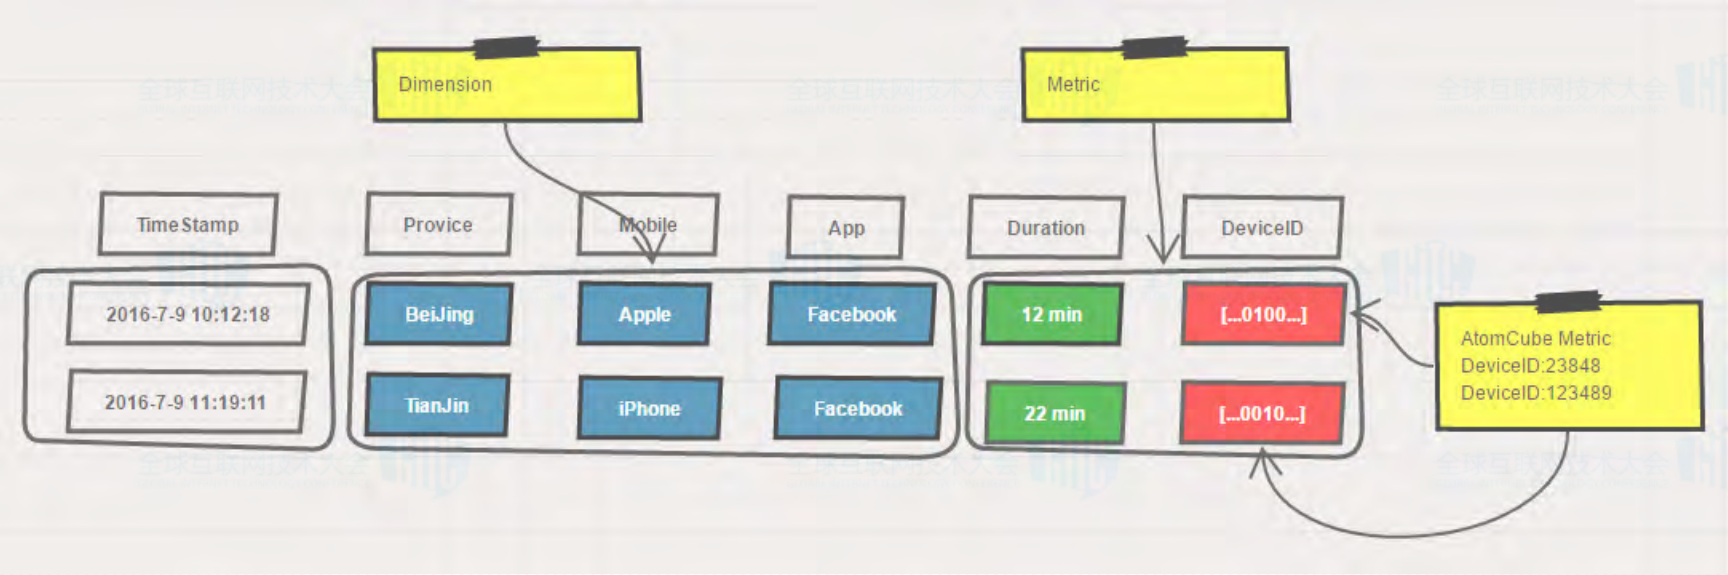

Main Points:

aggregation:

the atomcube naturally support.(UNION), so this feature can be applied in Rollup, QueryAggregation, and postAggregation processes.

Exactly cardinality:

the size of the atomcube after aggregation.

Join:

The atomcube in query result can do intersect if they have same means.

More:

Can do union and not operation

The Benefits:

- The huge bitmap stored separately and loaded in clustered nodes(Druid's historical).
- The computation can be accomplished in distributed environment (rollup, aggregation and historical) .
- Avoid Data skewed(Druid's load balance)

Nonintrusive

- Did not touch any existence Druid code.
- Just added an extension - druid-atomcube-0.9.0.jar
- Install
  - Put the druid-atomcube-0.9.0.jar under /druid-0.9.0/extensions/druid-atom-cube
  - add druid.extensions.loadList=["druid-atom-cube"] in common.runtime.properties file
  - startup all nodes

Implementation glance

- new DruidModule
- Aggregator and AggregatorFactory
  deserialize, metric, aggregate
- Query
  - Defined new query url:/druid/v2/atomcube
  - Defined new query structure.
  - Parallel running multiple queries on difference tables.

Schema Definition: define atomcube metric

```
"dataSchema":{
  "dataSource":"wikiticker",
  "granularitySpec":{
    ......
  },
  "parser":{
    ......
  },
  "metricsSpec":[
    {
      "name":"count",
      "type":"count"
    },
    {
      "name":"user_unique",
      "type":"hyperUnique",
      "fieldName":"hyperUser"
    },
    {
      "name":"user_atomcube",
      "type":"atomCube",
      "fieldName":"uuid"
    }
  ]
}
```

Query: new structure

```
{
 "queryType": "atomCube",
 ...
 "queries":   {
  "query1":{...}, // each query must include atomcube aggregation
  "query2":{...},
  "query3":{...},
 },
"postAggregations":[
 {
  "type": "atomCubeSet",
  "name": "test_set",
  "func": "INTERSECT",
  "fields": ["query1","query2","query3"]
 },
 {
  "type": "atomCubeSize",
  "name": "test_size",
  "field": "test_set"
 },
 {
  "type": "atomCubeRaw",
  "name": "test_raw",
  "format": "LIST",
  "field": "test_set"
 } ]
```

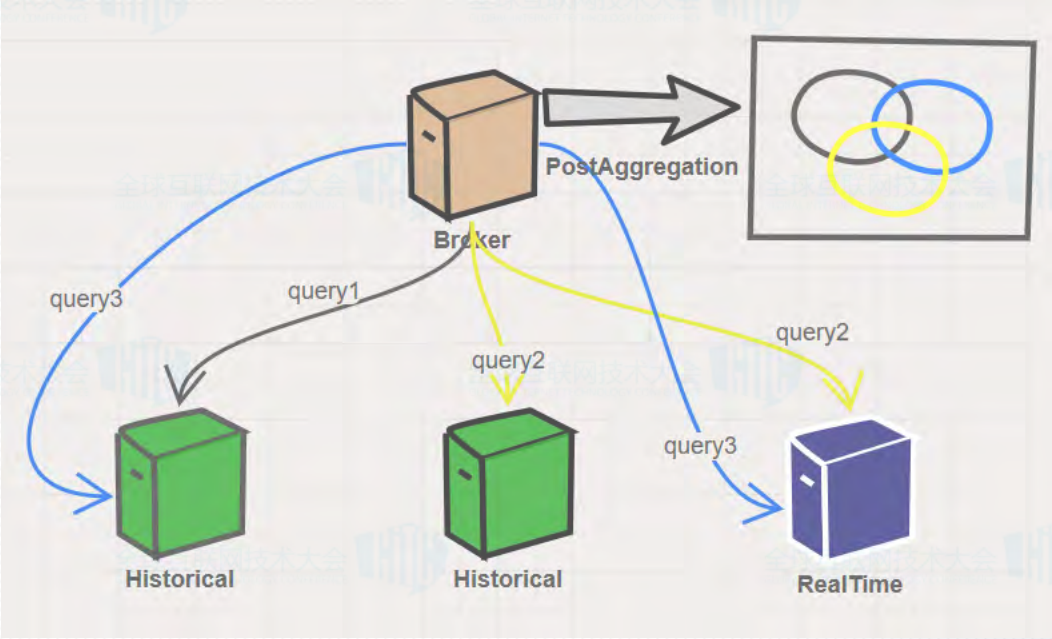

Query Result:

[{
  "test_raw" : [177, 1411, 2086, 2580, 3237, 3708, 3855, 6031, 6868, 8167, 8668, 9647, 10138],
  "test_size" : 13,
  "test_set" : "OjAAAAEAAAAAAAwAEAAAALEAgwUmCBQKpQx8Dg8PjxfUGucf3CGvJZon"
}]

Based on Calcite Druid Adapter

Refactor the Rules and QueryNode to Support Atomcube Query with standard SQL, like:

```
sql1 = "select distinct "uuid" from "wiki2" where "namespace  " = 'Wikipedia'";
sql2 = "select distinct "uuid" from "wiki1" where "countryName" = 'France'";
sql = "select distinct count(*) from (" + sql1 + " union " + sql2 + ")";
```

TalkingData

THANKS!

官网 / www.talkingdata.net

微博 / @TalkingData

微信 / TalkingData

服务支持 / support@tendcloud.com